1. What this app does

Brightspeed Call Analytics tags Billing and Payments call-center rows with a governed 104-leaf taxonomy. Upload monthly Excel workbooks, optionally filter by DNI and Billing/Payments scope, then run local embedding-based tagging. Review progress and download results from the Results tab.

2. Quick start

- Open the UI (default http://localhost:8501 or your VM IP:8501).
- Analyze calls: upload .xlsx files or pick existing datasets/.
- Set Billing / Payments and DNI toggles (see below).
- Choose Monthly, Custom, or Full dataset scope.
- Click Analyze and add tags, then open Results (auto-refresh is on by default).

3. Billing / Payments (BP) toggles

On Analyze calls, use the Billing calls and Payment calls toggles to control which Primary Call Reason rows are kept:

- Both on (default): keep Billing and Payments rows.
- Billing only: keep Billing rows.
- Payments only: keep Payments rows.
- Both off: not allowed - enable at least one before Analyze.

4. DNI filter toggle

Turn on Filter by DNI, then enter one or more DNIs (comma or newline separated). Matching is exact after trimming and ignoring case. Leave the toggle off to include all DNIs.

5. Other column filters

The Filters tab has optional advanced filters (BAN, Account Type, Network Type, Queue Category, etc.). Enable a column, then enter allowed values. Empty enabled filters block the run until you add values or uncheck the column.

6. Results, downloads, and cleanup

- Results shows status, metrics, and a live processing log while a run is active.
- When Complete, download tagged Excel and use pivot verification tables.
- Remove dataset files: delete Excel inputs you no longer need.
- Remove processed runs: delete finished ui_runs folders.

7. Docker (CPU / Google VM)

Seed test files then start: mkdir -p datasets && cp testdata/*.xlsx datasets/ && docker compose up --build -d. Open port 8501 in your cloud firewall. First tagging run may download the BGE embedding model into the hf_models volume.

8. Tips

- Prefer frozen prototypes in code/artifacts so tagging skips rebuild.
- Embedding batch size is 1000; give the VM several GB of RAM.
- Upload rejects duplicate filenames and workbooks missing required columns.